



CIS 455/555: Internet and Web Systems

Pastry; Message Queueing

March 13, 2013



Plan for today

- A few words about the team project 
- Pastry
 - Differences to Chord
 - API basics
- Message Queueing
- Remote Procedure Calls
 - Abstraction
 - Mechanism
 - Stub-code generation

Relevant
for HW3!

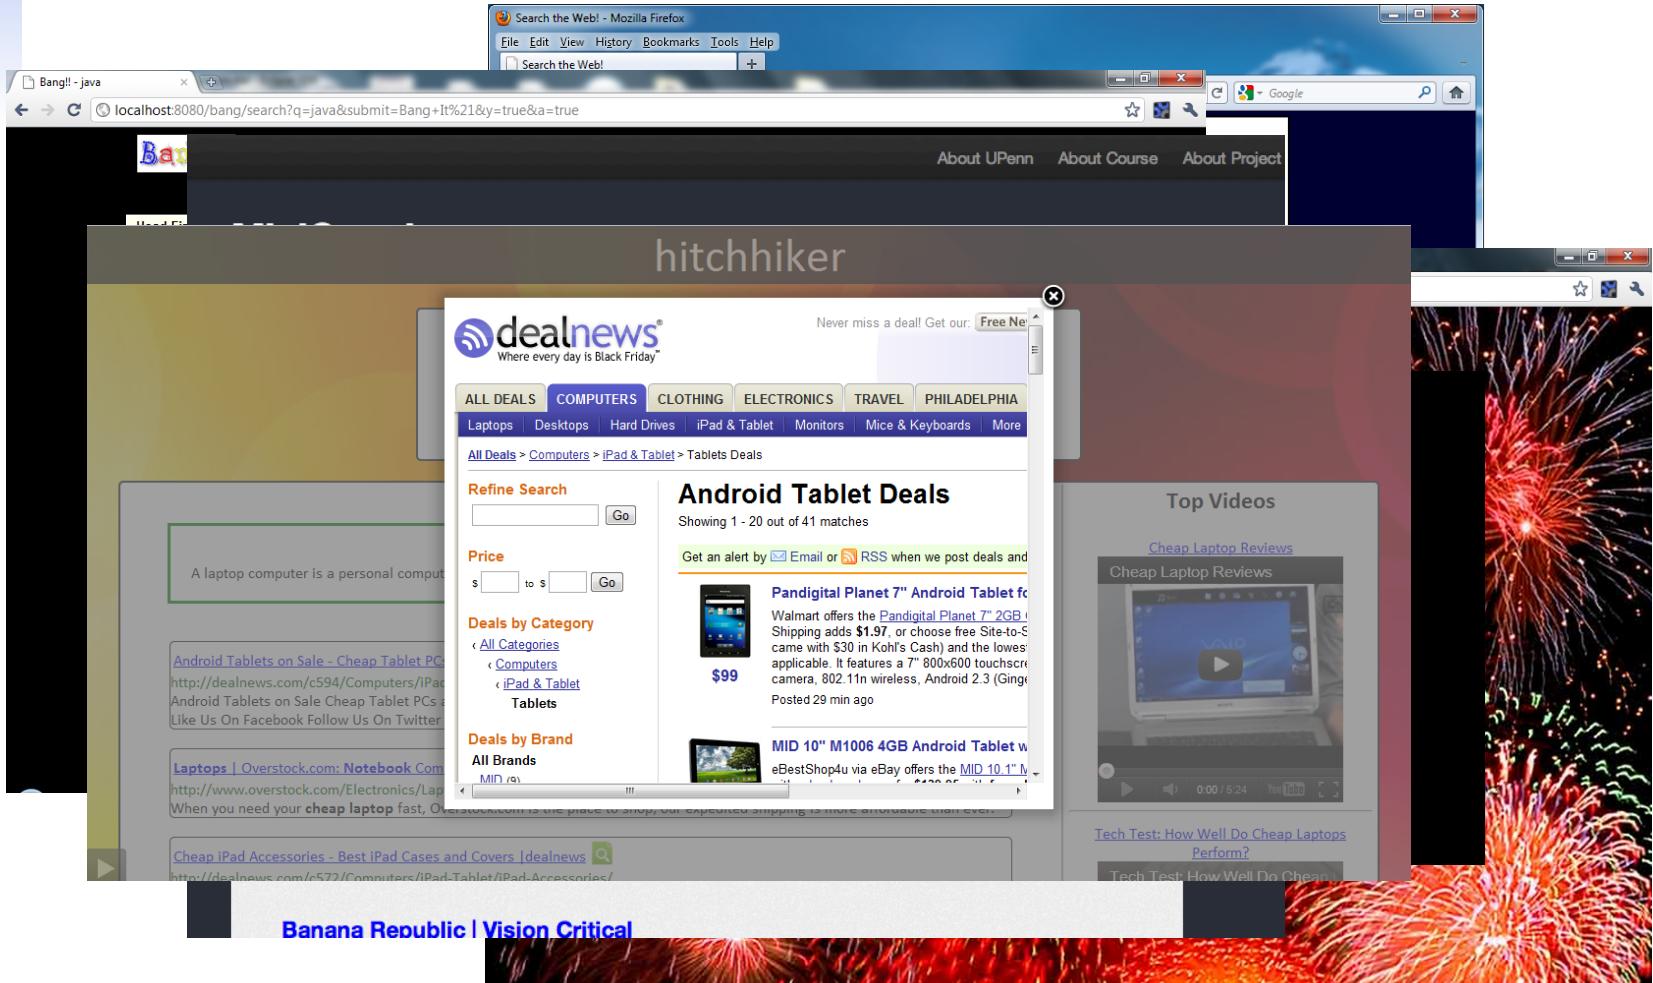


The team project

- Task: Build a P2P-based search engine
- Should consist of four components:
 - Crawler, based on your crawler (HW2) and Pastry (HW3)
 - Indexer, based on Pastry (HW3) and BerkeleyDB
 - PageRank, based on MapReduce
 - Search engine and user interface
 - Draft specs are available on the course web page
- Deploy & evaluate on Amazon EC2
 - Need to evaluate performance and write final report
 - Amazon has donated credit codes for this assignment
 - Will send out credit codes soon



Some of last year's projects





Logistics

- Rough timeline (preliminary):
 - Late March: Begin project planning
 - Early April: Initial project plan due
 - April 29: Official code submission deadline
 - April 30-May 3: Project demos
 - May 7: Final report due (hard deadline!)
- Todo: Form project groups
 - Each team should have 4 members
 - 5-member group requires approval & needs to do some extra credit tasks
 - One person from each group should send me a list of group members by Friday
 - I may have to split or merge some groups



The Google award



- The team with the best search engine will receive an award (sponsored by Google)
 - Criteria: Architecture/design, speed, reliability, quality of search results, user interface, written final report
 - Winning team gets four Android cell phones
 - Winners will be announced on the course web page



Some 'lessons learned' from last year

- The most common mistakes were:
 - **Started too late**; tried to do everything at the last minute
 - You need to leave enough time at the end to a) crawl a sufficiently large corpus, and b) tweak the ranking function to get good results
 - Underestimated amount of **integration work**
 - Suggestion: Define clean interfaces, build dummy components for testing, exchange code early and throughout the project
 - Performance issues
 - Do NOT use Pastry to transfer large amounts of data; use it for small, infrequent coordination messages!
 - Underestimated EC2 deployment
 - Try your code on EC2 as early as possible
 - Unbalanced team
 - You need to pick your teammates wisely, make sure everyone pulls their weight, keep everyone motivated, ...



Announcements

- HW2 MS2 is due on March 25th
 - Try to finish by Friday!
 - You can attend office hours / lab sessions
 - You'll have time for testing, and to fix unforeseen problems
 - Why not get started today?
- Second midterm is on April 22nd **at 4:30pm**
 - Likely location: DRL A1
- Reading for today:
 - A. Rowstron, P. Druschel: "Pastry: Scalable, distributed object location and routing for large-scale peer-to-peer systems"
 - <http://www.freepastry.org/PAST/pastry.pdf>



Plan for today

- A few words about the team project ✓
- Web services (continued) ← NEXT
- Pastry
 - Differences to Chord
 - API basics
- Message Queueing

Relevant
for HW3!



The Facebook API

- What you can do:
 - Read data from profiles and pages
 - Navigate the graph (e.g., via friends lists)
 - Issue queries (for posts, people, pages, ...)
 - Add or modify data (e.g., create new posts)
 - Get real-time updates, issue batch requests, ...

- How you can access it:
 - Graph API
 - FQL
 - Legacy REST API



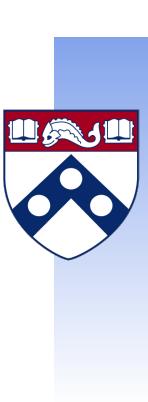
JSON

"Object": Unordered collection of key-value pairs

```
{  
    "firstName": "John",  
    "lastName": "Smith",  
    "age": 25,  
    "address": {  
        "streetAddress": "21 2nd Street",  
        "city": "New York",  
        "state": "NY",  
        "postalCode": 10021  
    },  
    "phoneNumber": [  
        { "type": "home", "number": "212 555-1234" },  
        { "type": "fax", "number": "646 555-4567" }  
    ]  
}
```

Array (ordered sequence of values; can be different types)

- Another standard for data interchange
 - "JavaScript Object Notation"; MIME type application/json
 - Basically legal JavaScript code; can be parsed with eval()
 - Caution: Security!
 - Often used in AJAX-style applications
 - Data types: Numbers, strings, booleans, arrays, "objects"



The Graph API

The screenshot shows the Facebook Graph API Explorer interface. On the left, there's a sidebar with a list of fields: Node: 1074724712, checked for id, age_range, locale, and location. In the center, a GET request is being made to https://graph.facebook.com/1074724712?fields=id,age_range,locale,location. On the right, the JSON response is displayed:

```
{
  "id": "1074724712",
  "age_range": {
    "min": 21
  },
  "locale": "en_US",
  "location": {
    "id": "101881036520836",
    "name": "Philadelphia, Pennsylvania"
  }
}
```

- Requests are mapped directly to HTTP:
 - [https://graph.facebook.com/\(identifier\)?fields=\(fieldList\)](https://graph.facebook.com/(identifier)?fields=(fieldList))
- Response is in JSON



The Graph API

- Uses several HTTP methods:
 - GET for reading
 - POST for adding or modifying
 - DELETE for removing
- IDs can be numeric or names
 - /1074724712 or /andreas.haeberlen
 - Pages also have IDs
- Authorization is via 'access tokens'
 - Opaque string; encodes specific permissions (access user location, but not interests, etc.)
 - Has an expiration date, so may need to be refreshed

Select Permissions

User Data Permissions Friends Data Permissions Extended Permissions

<input checked="" type="checkbox"/> email	<input type="checkbox"/> publish_actions	<input type="checkbox"/> user_about_me
<input type="checkbox"/> user_actions.music	<input type="checkbox"/> user_actions.news	<input type="checkbox"/> user_actions.video
<input type="checkbox"/> user_activities	<input type="checkbox"/> user_birthday	<input type="checkbox"/> user_education_history
<input type="checkbox"/> user_events	<input type="checkbox"/> user_games_activity	<input type="checkbox"/> user_groups
<input type="checkbox"/> user_hometown	<input type="checkbox"/> user_interests	<input type="checkbox"/> user_likes
<input type="checkbox"/> user_location	<input type="checkbox"/> user_notes	<input type="checkbox"/> user_photos
<input type="checkbox"/> user_questions	<input type="checkbox"/> user_relationship_details	<input type="checkbox"/> user_relationships
<input type="checkbox"/> user_religion_politics	<input type="checkbox"/> user_status	<input type="checkbox"/> user_subscriptions
<input type="checkbox"/> user_videos	<input type="checkbox"/> user_website	<input type="checkbox"/> user_work_history

Basic Permissions already included by default

Get Access Token **Cancel**



Other API options

- Facebook Query Language (FQL)
 - SQL-style queries over the data provided via the Graph API
 - Example: `SELECT name FROM user WHERE uid=me()`
 - Supports 'multi-queries' (JSON-encoded dictionary of queries)
 - 72 different tables are available for querying
- Legacy REST API
 - See description of REST in the previous lecture
 - In the process of being deprecated
 - Large number of methods available
 - Examples: `friends.get`, `comments.add`, `status.set`, `video.upload`, ...



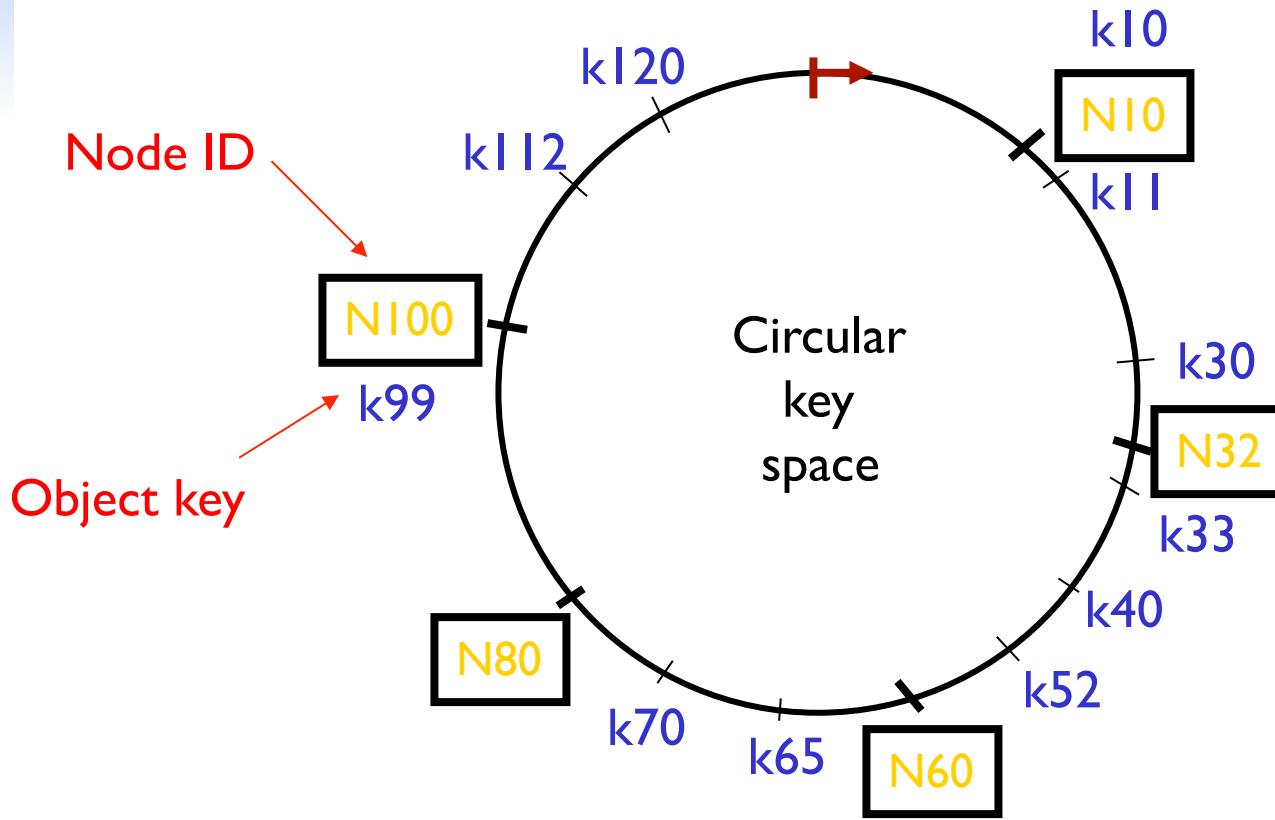
Plan for today

- A few words about the team project ✓
- Pastry ←
 - Differences to Chord
 - API basics
- Message Queueing

Relevant
for HW3!



Remember Chord?



- Large circular key space; objects and nodes have keys
- Key-based routing: route(message, key)



From Chord to Pastry

- What we saw was the basic data algorithms for the Chord system
- Pastry is a slightly different:
 - It uses a slightly different mapping mechanism (closest neighbor, not successor)
 - It has some extra features (leaf set, proximity-aware routing)
 - It allows for replication of data and finds the closest replica
 - It's written in Java, not C
 - ... And you'll be using it for HW3 and the final project!



FreePastry

- An open-source Java implementation of Pastry
 - Main web page:
<http://www.freepastry.org/>
 - Trac-Wiki:
<https://trac.freepastry.org/>
 - Tutorial:
<https://trac.freepastry.org/wiki/FreePastryTutorial>
 - Frequently Asked Questions:
<https://trac.freepastry.org/wiki/FreePastryFAQ>
 - Discussion group archives:
<https://mailman.rice.edu/pipermail/freepastry-discussion-l/>
 - Publications (if you're interested):
<http://www.freepastry.org/pubs.htm>



Pastry API Basics

- Keys/Node-IDs implement `rice.p2p.commonapi.Id`
 - Typical key space is $0..2^{160}-1$ (SHA-1)
 - Generated by factories (`NodeIdFactory`): `CertifiedNodeIdFactory`, `RandomNodeIdFactory`, `IPNodeIdFactory` - we'll use the latter
 - Each node has a `NodeId` (which implements `Id`) + `Ids` are used for routing
- Nodes are logical entities
 - A single physical machine can have more than one (virtual) node
 - Created by a `PastryNodeFactory` - in our case, `SocketPastryNodeFactory` (but there are others, e.g., for local simulations)
- All Pastry nodes have built in functionality to manage routing
 - Derive from “common API” `rice.p2p.commonapi.Application`



Creating a P2P Network

```
public DistTutorial(int bindport, InetSocketAddress bootaddress, Environment env) {  
    NodeIdFactory nidFactory = new RandomNodeIdFactory(env);  
    PastryNodeFactory factory = new SocketPastryNodeFactory(nidFactory, bindport, env);  
    PastryNode node = factory.newNode();  
    /* register your applications here */  
    node.boot(bootaddress);  
    synchronized(node) {  
        while(!node.isReady() && !node.joinFailed()) {  
            // delay so we don't busy-wait node.wait(500); abort if can't join  
            if (node.joinFailed())  
                throw new IOException("Could not join the FreePastry ring. "+  
                    "Reason:"+node.joinFailedReason());  
        }  
    }  
    System.out.println("Finished creating new node "+node);  
}
```

- Example code in `DistTutorial.java`
 - Creates a new node, which joins the overlay ("ring")
 - Need to provide address of an existing member of the ring ("bootstrap node"); if no ring exists yet, it starts a new one
- No need to call a simulator – this is real!
 - For more information, see FreePastry tutorial:
https://trac.freepastry.org/wiki/tut_lesson3



Pastry API basics

- Based on a model of routing **messages**
 - Derive your message from rice.p2p.commonapi.Message
 - Every message gets an **Id** corresponding to its key
- Concept of **node handles** (**NodeHandle** class)
 - Can be used to talk directly to a specific node (why would you need this?)
 - Internally, has a NodeId, and IP address, and a port number
- Concept of **endpoints** (**Endpoint** class)
 - You write an application (rice.p2p.commonapi.Application) and register it with the Pastry node to get one
 - Nodes can have multiple endpoints; provide string as identifier



Routing API in Pastry

- Call `endpoint.route(id, msg, hint)` to send a message
 - If id is given and hint is null, use key-based routing
 - If hint (a NodeHandle) is given and id is null, send message directly to node
 - If both are given, key-based routing will be used, and hint will be first hop
 - Not reliable -- can lose messages (what does this mean for your app?)
- At each intermediate point, Pastry makes an upcall (**forward**) to the corresponding application
 - You are given the message that is being routed
 - You can read out (or change) the key, the contents, and the NodeHandle of the next hop
- At the end, Pastry makes a final upcall (**deliver**) to your application
 - Called on the node whose NodeId is closest to the message's Id
 - Is given the message and the Id



Transferring large amounts of data

- FreePastry uses a single TCP connection for its message-based traffic
 - If you route huge messages, this will stall FreePastry's maintenance traffic (and, obviously, other messages)
 - Likely result: Poor performance
- To transfer large amounts of data, use the **application-level socket interface**
 - Sender can open a direct connection to the receiver, without intermediate hops (requires a NodeHandle, though)
 - https://trac.freepastry.org/wiki/tut_app_sockets
- Use messages only for coordination
 - Don't send too many; "congestion control" may be needed



Making IDs

- Pastry has mechanisms for creating node IDs itself
- Obviously, we need to be able to create IDs for keys
- Need to use `java.security.MessageDigest`:

```
MessageDigest md = MessageDigest.getInstance("SHA");
byte[] content = myString.getBytes();
md.update(content);
byte shaDigest[] = md.digest();
```

```
rice.pastry.Id keyId = new rice.pastry.Id(shaDigest);
```



Creating a DHT abstraction with Pastry

We want the following:

- put (key, value)
 - remove (key)
 - valueSet = get (key)
-
- How can we use Pastry to do this?



Recap: Pastry

- A substrate for decentralized systems
 - Implements key-based routing
 - Similar to Chord, but has some additional features, e.g., leaf sets and proximity neighbor selection
 - Open-source Java implementation available (FreePastry)
- Main FreePastry abstractions:
 - Id and NodeId (keys from a large key space)
 - Endpoint (for sending/receiving messages)
 - Node (has a NodeId, provides endpoints, handles routing)
 - Application (superclass for your own KBR application)



Plan for today

- A few words about the team project ✓
- Pastry ✓
 - Differences to Chord ✓
 - API basics ✓
- Message Queueing ← NEXT

Relevant
for HW3!



Why message queueing?

- So far, we have seen synchronous/transient communication
 - Example: REST call to a web service
- But what if applications are more loosely coupled?
 - Maybe the service we're interacting with isn't always available (batch process etc.)

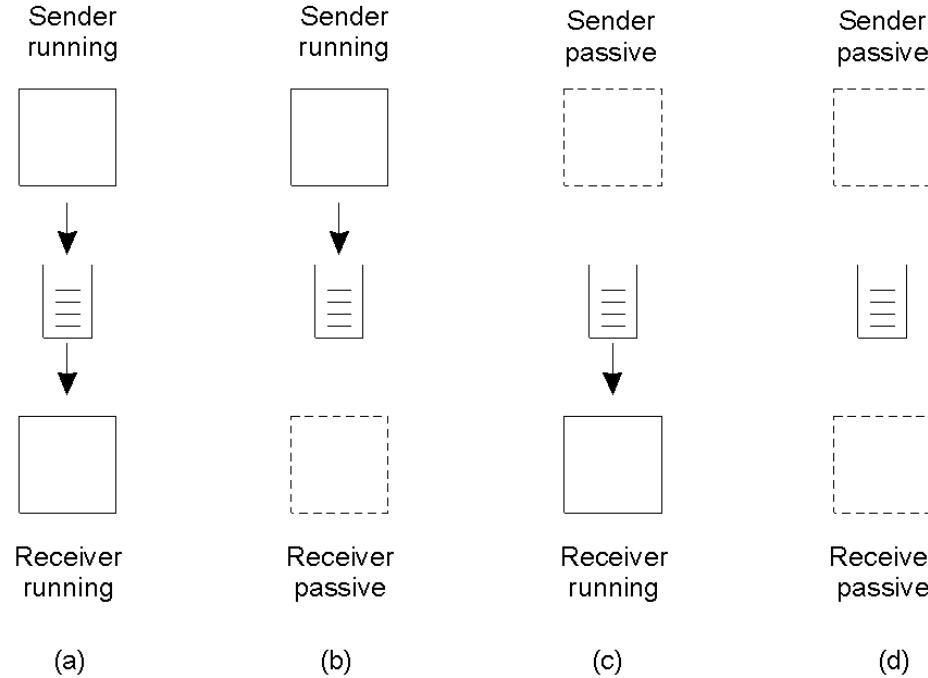


Types of communication

- There are several types of communication:
 - Persistent vs transient
 - If the sender and the receiver are not running, is the message lost?
 - Synchronous vs asynchronous
 - Does the sender continue immediately, without knowing that the message has been accepted?
- What would be an example of...
 - Synchronous, transient communication?
 - Synchronous, persistent communication?
 - Asynchronous, transient communication?
 - Asynchronous, persistent communication? ← Need queues!
Hence MQM
 - Which of these are tightly/loosely coupled?



Message-Queuing Model



- Apps communicate by inserting messages into queues
 - Sender only knows that message has been inserted into queue, not that it has been seen or processed by the recipient (real-world analogies?)
 - Four combinations; see above
- Implemented e.g., by IBM WebSphere MQ, Oracle AQ, ...
 - Similar: Amazon SQS



Message-Queuing API

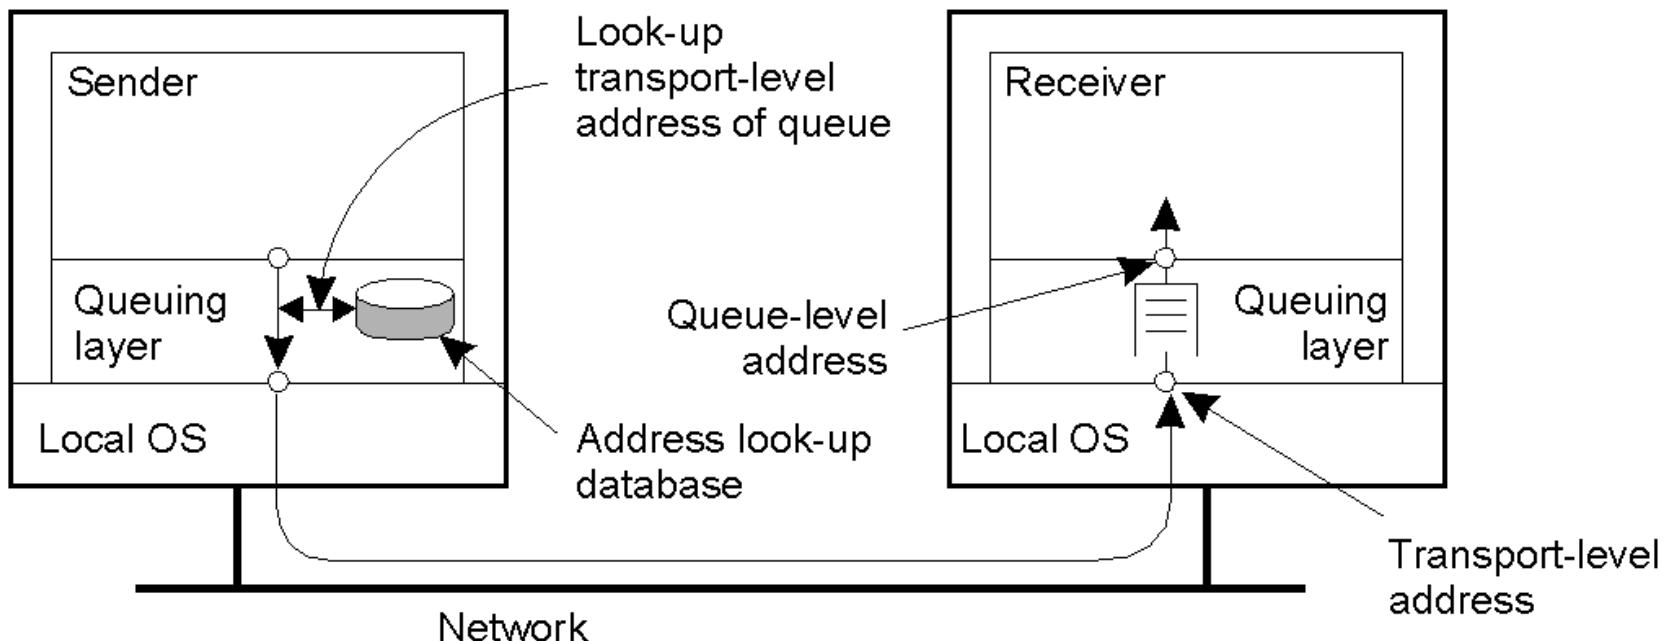
- Basic interface to a queue in a message-queuing system.

Primitive	Meaning
Put	Append a message to a specified queue
Get	Block until the specified queue is nonempty, and remove the first message
Poll	Check a specified queue for messages, and remove the first. Never block.
Notify	Install a handler to be called when a message is put into the specified queue.



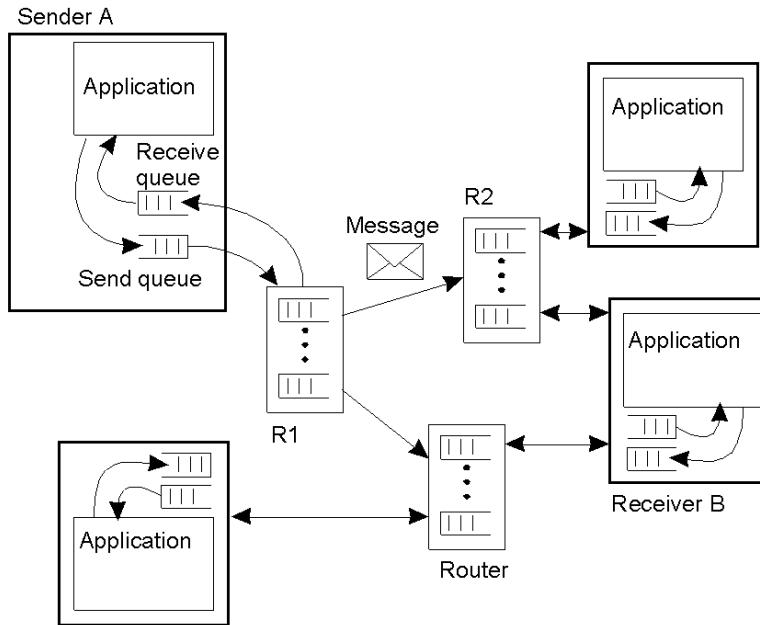
General Architecture of a MQ System

- The relationship between queue-level addressing and network-level addressing.





MQ System with routers



- Problem: Each queue manager has to be able to find the location of any given queue
 - Complexity / scalability challenges
- Solution #1: Directory service
- Solution #2: Relays / routers



Benefits of message queueing

- Allows both synchronous (blocking) and asynchronous (polling or event-driven) communication
- Ensures messages are delivered (or at least readable) in the order received
- The basis of many transactional systems
 - e.g., Microsoft Message Queue (MMQ), IBM MQseries, etc.